

Scope of You

Aura Yavary

Learning Where Personalization Is Allowed to Matter

THE QUESTION

A model can know a user's preference and still apply it in the wrong place. Scope of You asks where a known preference is allowed to change behavior, using three typed obligations: APPLY, SUPPRESS, and MUST-NOT-AFFECT.

THE FALSIFICATION

The first contract optimizer looked strong on an explicit, co-designed test and failed on an independently worded implicit benchmark. The benchmark was kept; the training intervention changed. The failed result remains in the release.

STRONGEST EXECUTED EVIDENCE

ContractBench-Implicit: 1,460 cases, 20 held-out users. PCO-Robust = 0.7912 vs BATPO = 0.6866 and SFT = 0.7135. Hierarchical user-to-case bootstrap vs BATPO: +0.1046, 95% CI [+0.0865, +0.1224], positive on 20/20 held-out users. Five-seed mean = 0.7910 +/- 0.0008.

THE TRADEOFF

This is not a universal win. PersonalBench utility drops to 0.9448 from 0.9905 for SFT; the predeclared 2-point utility non-inferiority gate fails. The project reports the selectivity-utility frontier rather than collapsing it into a single score.

WHY IT IS MORE THAN A BENCHMARK DEMO

Matched counterfactual interventions; retained negative result; hierarchical statistics; non-inferiority testing; protected-dimension leakage analysis; a stateful permission/recovery evaluator with stale-memory and untrusted-tool perturbations; a frozen 36-run two-family LLM matrix; blinded human-eval protocol; machine-checked claim/evidence and experiment provenance.

WHAT IS NOT PROVEN

No frontier-LLM SOTA, human-preference superiority, calibrated uncertainty, production-agent safety, or longitudinal real-user result is claimed. These remain explicit external gates.

NEXT DECISIVE TEST

Execute the frozen 36-run study on two independent 7-8B model families, run public personalization/selectivity benchmarks, preserve the 2-point utility gate, and collect blinded human scope/appropriateness judgments. If the typed-scope advantage survives, the claim graduates from controlled mechanism evidence to an LLM-scale result.